

Range-Filtering Approximate Nearest Neighbor Search via First-Split Range Decomposition and Suffix-Prefix Indexing [Technical Report]

Abstract—Given a query vector and n input data objects, range-filtering approximate nearest neighbor search (RFANNS) aims to find its nearest neighbors whose numeric attributes fall within a query range, supporting various applications in databases, information retrieval, and machine learning. Existing methods for RFANNS either consult more than a constant number of indexes per query or build an index over a range evidently larger than the query range. The two drawbacks impair their scalability to large-scale queries.

To improve search efficiency, we propose **FRASP**, a novel method that employs a first-split range strategy with suffix and prefix HNSWs in a segment tree. **FRASP** answers any query with only three segment-specific graph indexes and evaluates edges per object with running time linear with the maximum degree in graph indexes. The superior performance of **FRASP** is achieved by four key designs. (i) We apply the first-split position in the segment tree to split the query range into a suffix and a prefix range, contained by the left and right child segments of a parent segment respectively. (ii) During index construction, we build suffix HNSWs for left child segments and prefix HNSWs for right child segments, enabling arbitrary query ranges to be supported by two one-sided ranges. (iii) A unified search is performed over the suffix, prefix, and parent segments, to explore objects within the query range without redundant computations. (iv) We design succinct prefix/suffix array indexes to locate relevant HNSW layers and initial entry points efficiently via simple array lookup. Extensive experiments indicate that **FRASP** outperforms state-of-the-art methods in terms of both query efficiency and result quality.

Index Terms—approximate nearest neighbor search, range filtering, graph index, HNSW, vector database

I. INTRODUCTION

In the era of machine learning and big data, many real-world entities are represented as vectors to support downstream tasks. Efficient management of high-dimensional vector data becomes an essential technique in many applications, including recommendation systems [1], information retrieval [2], machine learning systems [3], [4], language models [5], and database systems [6]. A fundamental task in vector processing is nearest neighbor search (NNS), which aims to find the closest vectors to a query from a large dataset based on a distance metric such as Euclidean distance.

With the rapid growth of vector dimension and dataset size, exact nearest neighbor search (NNS) becomes computationally expensive and incurs high latency. To improve efficiency, the distance requirement is relaxed to approximate NNS (ANNS), which finds vectors that are approximately nearest to the query. This trade-off significantly improves query efficiency while maintaining high result quality. Various indexes have

been developed to support ANNS, many of which are based on Hierarchical Navigable Small World (HNSW) graphs [7], achieving high recall and throughput [8]–[10].

In industrial scenarios and vector databases [11], [12], data objects often include numerical attributes (e.g., timestamps or prices) alongside vectors. Many applications require joint querying of vectors and their attributes. For example, on e-commerce platforms, a user can upload a query product image and specify a price range to find similar products within the range. On video platforms like YouTube, video recommendations based on a query vector (derived from a user’s watch history) must comply with age-related regulations by applying an age filter. These queries are known as range-filtering ANNS (RFANNS), aiming to find the nearest neighbors of a query vector whose attributes fall within the specified range filter (i.e., *in-range*), while excluding *out-of-range* objects.

A profusion of methods are proposed to solve RFANNS. Simple adaptations to previous methods (e.g., pre-filtering, post-filtering, and Milvus [11]) for ANNS are often inefficient, since they are not especially designed for RFANNS. Recent advancements [9], [10], [13], [14] have developed efficient indexes and query techniques dedicated for RFANNS. Even so, existing methods still need to be improved. For SeRF [10] and WoW [14], their indexes contain many out-of-range or invalid edges, which occupy extra space and should be skipped during search. UNIFY [15] and iRangeGraph [9] maintain too many fragmented indexes, resulting in evident redundancy in indexes or poor connectivity between ranges. Moreover, UNIFY and iRangeGraph need to consult many indexes, leading to higher query time. To elaborate, UNIFY may examine $O(S)$ segment-wise chunks in hybrid filtering, with $S = 8$ in its default configuration, and iRangeGraph consults $O(\log n)$ graph indexes for each query. Thus, more efficient approaches should be designed to accelerate RFANNS.

In this paper, we propose an efficient method for RFANNS called **FRASP**¹, based on a segment tree which takes each segment as a node and each segment is recursively split at its midpoint. To achieve high efficiency, **FRASP** involves four techniques tailored for RFANNS. (1) **The first-split strategy**. Given a query range $[L, R]$, we first find from the segment tree the largest segment $[l, r]$ whose midpoint (denoted as fs) is in $[L, R]$. Then, the query range is split into two ranges $[L, fs]$ and $(fs, R]$, called *suffix range* and *prefix range* resp, which

¹First-split **R**ange exploration using **S**uffix and **P**refix HNSWs

are contained by the children of $[l, r]$ (i.e., $[l, fs]$ and $[fs, r]$ resp.). In this way, we only need to focus on three segments $[l, r]$, $[l, fs]$, and $[fs, r]$ for a query. (2) **Suffix/prefix HNSW**. For efficient search in any segment, we build an HNSW for each segment, whose role in the segment tree is a left child or a right child. If a segment is a left child, we construct a suffix HNSW for it. Otherwise, we build a prefix HNSW for the segment. By default, the root segment maintains a prefix HNSW. (3) **Unified search**. To leverage the inter-edges between child segments, we utilize the HNSWs of not only the child segments but also the parent segment $[l, r]$. On this basis, we design a unified search process over the parent segment and the two child segments, to share the intermediate results and a global lower bound for pruning and early termination. This unified search ensures exploration is restricted to in-range objects, effectively simulating a single HNSW search for the query range $[L, R]$, avoiding redundant computations. (4) **Succinct Array Indexes**. Note that, HNSW is not originally designed for RFANNS. Directly searching from the top-most layer leads to unnecessary out-of-range computation. Thus, we design succinct array indexes that enable efficient identification of the entry points through simple array lookup. With the above four techniques, FRASP processes each query using only three segments and their HNSWs. Each processed object appears in at most two of these HNSWs, so at most $2M_0$ ordinary outgoing edges are accessed per object, where $M_0 = 2M$ is the maximum base-layer degree of HNSW. Thus, the expected query processing time per object is $O(M)$.

Extensive experiments on various benchmark datasets demonstrate that FRASP achieves significant improvements over existing methods, in terms of both efficiency and quality. In summary, our contributions are presented as follows.

- We propose the FRASP method to solve RFANNS efficiently. For each query, FRASP only consults three segments and their HNSWs, much less than existing methods.
- We design a first-split strategy to process query ranges in the segment tree. This strategy is not tied to HNSW and can be instantiated with ANNS indexes that support attribute-ordered insertion.
- Accordingly, we design suffix and prefix HNSWs for the query ranges. Moreover, we design succinct prefix/suffix arrays to efficiently identify the entry points of HNSWs.
- We design a unified search process over the three HNSWs, which maintains the connectivity between child segments and ensures efficient exploration within the query range.
- Experiments validate that FRASP outperforms state-of-the-art methods in terms of query efficiency and result quality.

II. PRELIMINARIES

A. Problem Definition

Given a set of n vectors $\mathcal{D} = \{o_1, o_2, \dots, o_n\}$ and a query vector q in $\mathbb{R}^d$, the nearest neighbor search (NNS) problem aims to find the top k nearest neighbors in $\mathcal{D}$ under certain distance metric δ , e.g., the Euclidean distance [16]. In many applications, each object $o \in \mathcal{D}$ is associated with a numeric

attribute $o.a$ in addition to the vector data, such as size, time, or price. Accordingly, the aim becomes finding the nearest neighbors within a specific range $[a_L, a_R]$ of attribute values, which is formulated as the range-filtering NNS (RFNNS) problem in Definition 1.

Definition 1 (RFNNS [9], [10]). *Given a set $\mathcal{D}$ of objects in $\mathbb{R}^d$ where each object o is associated with an attribute $o.a$, a query vector $q \in \mathbb{R}^d$, and a query range $[a_L, a_R]$, the RFNNS query is to find a set Φ^* of k objects from $\mathcal{D}$ such that:*

- 1) $\forall o_i \in \Phi^* : o_i.a \in [a_L, a_R]$, and
- 2) $\forall o_i \in \Phi^*, \forall o_j \in \mathcal{D} \setminus \Phi^* : o_j.a \in [a_L, a_R] \Rightarrow \delta(q, o_i) \leq \delta(q, o_j)$.

Exact results are computationally expensive to compute, and thus *approximate RFNNS* (RFANNS) algorithms are often studied, aiming to efficiently find an approximate top- k result Φ that closely approximates the true nearest neighbors Φ^* within the query range. The performance is often measured by recall $|\Phi \cap \Phi^*|/|\Phi^*|$.

TABLE I
FREQUENTLY USED NOTATIONS

Notation	Description
n	Number of objects in the dataset
$q, [L, R]$	Query vector in $\mathbb{R}^d$ and Query range
fs	First-split position in the segment tree for a query range
$\mathcal{D}$	Set of n objects in $\mathbb{R}^d$
$o.a$	Attribute value of object o
$[l, r]$	Segment in the segment tree
$\mathcal{H}_{\text{sfx}}$	Suffix HNSW for a left child segment
$\mathcal{H}_{\text{prf}}$	Prefix HNSW for a right child segment
$\mathcal{H}_p$	HNSW for the parent segment
$N_b[o_i]$	Outgoing neighbors of o_i at the b -th layer
$N_b^{\text{del}}[o_i]$	Pruned (deleted) neighbors of o_i at the b -th layer
M	Maximum degree parameter in HNSW
ef, ef_{con}	Candidate list size for search/construction
A	Array of highest HNSW layers per object
Z	Suffix-Max (or Prefix-Max) array depending on the segment
$\mathcal{L}_a$	The a -th layer in HNSW
$\mathcal{L}_0$	Base layer in HNSW
$\mathcal{L}_{\text{qmax}}$	Topmost query-relevant HNSW layer or layer id
ep_{qmax}	Query-relevant entry point in $\mathcal{L}_{\text{qmax}}$
$\mathcal{C}$	Candidate priority queue in search
$\mathcal{R}$	Result priority queue in search
$\delta(\cdot, \cdot)$	Distance function (e.g., Euclidean)

An object is referred to as *in-range* if its attribute is within the query range $[a_L, a_R]$, and as *out-of-range* otherwise. Following [9], [10], we assume that the objects in $\mathcal{D} = \{o_1, o_2, \dots, o_n\}$ are sorted in ascending order of their attribute values. If not, the objects can be pre-sorted based on their attribute values. This sorting establishes a one-to-one mapping between the index of an object and its attribute value. That is, for any $o_i, o_j \in \mathcal{D}$, if $i < j$, then $o_i.a \leq o_j.a$. This allows the index of an object to represent its attribute value. Given a query range $[a_L, a_R]$, it is easy to determine the smallest index L in $\mathcal{D}$ such that $o_L.a \geq a_L$, and the largest index R such that $o_R.a \leq a_R$. Thus, the problem is reduced to finding the k nearest neighbors of q among the objects in the index range $[L, R]$ of $\mathcal{D}$.

Table I summarizes the frequently used notations in this paper. Due to space constraints, the proofs of lemmas and theorems are deferred to our technical report [17].

B. The HNSW Index

Graph-based indexes (such as NSW [18], NSG [19], and Vamana [20]) exhibit superior performance for ANNS [7], [8], [21]. Their main idea is to build a graph called *proximity graph* where each vertex represents a vector and edges are established based on the distance between vectors. Based on NSW, HNSW organizes multiple layers in a hierarchical structure, where each layer is an NSW graph in essence. Given an HNSW $\mathcal{H}$, the layers are denoted as $\mathcal{L}_0, \dots, \mathcal{L}_{\max}$, where $\mathcal{L}_0$ is the base layer and $\mathcal{L}_{\max}$ is the topmost layer. The base layer $\mathcal{L}_0$ contains all vectors, while each upper layer contains an exponentially decreasing number of vectors. In each layer $\mathcal{L}_b$, a directed edge from o_i to o_j indicates that o_j is stored as a neighbor of o_i ; in this case, o_j is called an out-neighbor of o_i . The set of out-neighbors of o_i is denoted as $N_b[o_i]$, the cardinality of which is called the out-degree of o_i .

Algorithm 1: InsertObject(o)

Input: Entry point $ep_{\max}$ in $\mathcal{L}_{\max}$, ef_{con} , maximum connections M , the object o to be inserted
Output: Updated HNSW $\mathcal{H}$

```

1  $\mathcal{L}_a \leftarrow \text{SampleLevelRandomly}(M)$  // Sample
  a level for  $o$ 
2  $ep \leftarrow ep_{\max}$ ;
3 for  $\mathcal{L}_b$  from  $\mathcal{L}_{\max}$  to  $\mathcal{L}_{a+1}$  do
4    $ep \leftarrow \text{SearchLayer}(o, ep, ef = 1, \mathcal{L}_b)$ ;
5 for  $\mathcal{L}_b$  from  $\mathcal{L}_a$  to  $\mathcal{L}_0$  do
6    $\mathcal{W} \leftarrow \text{SearchLayer}(o, ep, ef_{\text{con}}, \mathcal{L}_b)$ ;
7    $N_b[o] \leftarrow$ 
     Select at most  $M$  nearest neighbors from  $\mathcal{W}$ ;
8   foreach  $v \in N_b[o]$  do
9      $N_b[v] \leftarrow \text{PruneExistingEdges}$ 
       ( $N_b[v] \cup \{o\}$ );
10 return Updated  $\mathcal{H}$ ;
```

HNSW Construction. An HNSW can be constructed by constantly inserting objects (e.g., o) with the InsertObject method in Algorithm 1. To begin construction, we set the maximum out-degree M of each vertex and the maximum number of candidates ef_{con} proposed as the nearest neighbors of o . Next, we sample a random level $\mathcal{L}_a$, downwards from which the object o will be inserted. Before that, starting from $\mathcal{L}_{\max}$ (Lines 3–4), we find the entry point for each layer above $\mathcal{L}_a$ by applying a SearchLayer procedure with $ef = 1$, which narrows the subsequent search space by providing a good nearest neighbor as the entry point. The details of SearchLayer are omitted and can be found in [7]. For each layer $\mathcal{L}_b$ not above $\mathcal{L}_a$, we search for ef_{con} nearest neighbors of o , denoted as $\mathcal{W}$ (Line 6), from which at most M neighbors are selected to form the out-neighbors $N_b[o]$ (Line 7). For each out-neighbor $v \in N_b[o]$, we create an edge between o and v . Moreover, in Lines 8–9, the existing edges in $N_b[v]$ are updated and pruned to maintain search efficiency. The procedure PruneExistingEdges uses the strategy of SelectNeighbors in [7]. Finally, the updated HNSW $\mathcal{H}$

is returned (Line 10). The expected number of edges created in HNSW is $O(nM(1+1/\ln(M)))$, and the term $1+1/\ln(M)$ is the expected number of layers an object appears in HNSW. M is usually 16–64, and thus the space cost of HNSW is $O(nM)$.

Search on HNSW. Given a query vector q , the search process in HNSW starts from $\mathcal{L}_{\max}$ and proceeds downwards to the base layer $\mathcal{L}_0$. In each layer, HNSW applies the SearchLayer procedure to find the nearest neighbor of o , which serves as the entry point for the next layer. In the base layer $\mathcal{L}_0$, SearchLayer is called again to find ef nearest neighbors of q , from which the top- k are returned as the final result.

III. RELATED WORKS

In this section, we first introduce approaches designed for general ANNS and then discuss methods on attribute filtering ANNS. Finally, we elaborate the works on RFANNS, which are most related to our work.

A. General ANNS

The research on ANNS can be roughly categorized into the methods based on the tree structure, hashing, graphs, and quantization. The tree-based methods [22], [23] leverage hierarchical structures to separate data, enabling fast searches with approximate pruning. However, they are less effective in high dimensions. Hashing-based methods [24]–[26] map similar points to the same buckets by probabilistic hash functions, providing scalability for large datasets. Graph-based methods [21], [27]–[29], such as HNSW [7] and DiskANN [20], build proximity graphs for highly accurate and tunable searches. Quantization-based methods [30]–[33] compress data into compact codes to reduce memory usage and accelerate distance computations. Some works also apply hybrid ideas to solve ANNS, e.g., [27]. For a more detailed tutorial, the readers may refer to [34], [35].

B. Attribute Filtering ANNS

This formulation extends ANNS by requiring returned objects to satisfy additional attribute predicates. The Native Hybrid Query framework is proposed in [36], which achieves up to $315\times$ speedup. Gollapudi et al., propose Filtered-DiskANN to support categorical and numerical filters under the disk-based setting [37]. The HQI system [38] is designed for ANNS in knowledge graphs. In [39], the authors propose a reconfigurable index for multimedia datasets with significant speedups. An index called ARKGraph is proposed in [40] to support range queries on vectors with numerical attributes, which is closely related to RFANNS. ACORN increases the number of edges in HNSW to handle low-selectivity queries with general predicates. Milvus supports attribute filtering through hybrid queries that combine vector similarity with metadata constraints [11].

C. Range Filtering ANNS (RFANNS)

Early works solve RFANNS by combining searching and filtering. For example, pre-filtering uses binary search to find

in-range objects and then computes distances. It is efficient only for very small query ranges (e.g., ≤ 1024 objects). Post-filtering first retrieves candidate nearest neighbors via ANNS, then applies the range filter to obtain k in-range neighbors. It becomes inefficient for selective queries with few in-range items. SuperPostfiltering [41] improves post-filtering by searching the shortest pre-built range that covers the query range. The covering range may be larger than the query range, causing extra distance computations on out-of-range objects.

Later works try to apply more complex indexes (such as HNSW, segment trees) to improve the efficiency. iRangeGraph [9] builds a segment tree over attribute values, attaching a graph index to each segment. For a query range, it consults $O(\log n)$ segments and constructs a dedicated graph on the fly. To process each object, the amortized time of iRangeGraph is $O(M + \log n)$, and the total space is $O(nM \log n)$. 2DSegmentGraph [10] extends HNSW by attaching attribute-related tags to each edge. For every edge, four tags (l, r, b, e) are stored to check query range bounds. An edge is eligible if the query range falls into these intervals. This approach may create many edges per node, leading to $O(n^2 M)$ worst-case space and an average space of $O(nM \log n)$. UNIFY [15] builds a unified HSIG index for RFANNS over one scalar attribute. It combines segment-wise graph structures, skip-list connections, and compressed global HNSW edges to support hybrid, pre-, and post-filtering within a single index, with range-aware strategy selection. The most recent method called WoW [14] sorts scalar attributes with a weighted balanced tree for window boundaries, and maintain hierarchical window graphs. In each layer, neighbors are restricted to a symmetric attribute window. It selects a landing layer to balance filter checks and distance computations.

IV. INDEX STRUCTURE AND CONSTRUCTION

An efficient RFANNS index should satisfy two requirements: search should be restricted to the query range, and the underlying proximity graph should remain navigable for high-recall search. A direct idea to solve RFANNS is partitioning the search space into many small segments and building indexes for them. This idea preserves range precision but impairs graph connectivity. Conversely, searching a larger segment can preserve connectivity but may incur unnecessary exploration.

To address the above dilemma, FRASP proposes to represent every query range with a constant number of segment indexes at its first-split level. Specifically, the index of FRASP is designed under four principles. (i) FRASP uses a first-split range representation to express the query range as two one-sided child subranges contained by a parent segment. Next, we construct an HNSW for each segment. As a result, each query only accesses three HNSWs. (ii) To construct the HNSWs for the two one-sided segments, FRASP inserts objects in attribute order and retains pruned edges with deletion tags. In this way, we can recover the HNSW that would have been constructed directly over the queried one-sided segment. (iii) The HNSW of a segment remains unchanged regardless of whether the

segment serves as a child or parent, without materializing additional bridge indexes. Then, FRASP performs a unified search to traverse the three HNSWs with shared candidate and result queues. (iv) To support range-aware HNSW initialization, FRASP maintains lightweight auxiliary arrays for each segment HNSW, so the search can start from an in-range upper-layer entry point without scanning layers.

A. First-Split Range Strategy

We adopt the segment tree to handle the ranges, which is a classic binary tree structure over the numerical range $[1, n]$. Each node represents a segment and is recursively split at its midpoint. Figure 1 shows an example of a segment tree over $[1, 16]$. Although many segments in the tree may intersect a query range, FRASP does not decompose the range into all canonical nodes. Instead, it identifies the first level of the segment tree at which the query range crosses a midpoint when traversing from the root. Given a query range $[L, R]$, let $[l, r]$ be the *highest* segment-tree node whose midpoint $mid = \lfloor (l+r)/2 \rfloor$ splits the query range, i.e., $L \leq mid < R$. Then, we call $[l, r]$ the *parent segment* of the query range, $fs = mid$ the *first-split position*, and the level that contains $[l, r]$ the *first-split level*. The two children of the parent segment, namely $[l, fs]$ and $[fs+1, r]$, are called the *left child segment* and the *right child segment*, respectively. In analogy, the HNSW built on $[l, r]$ will be referred to as the *parent HNSW*, while the suffix HNSW of $[l, fs]$ and the prefix HNSW of $[fs+1, r]$ will be referred to as the *child HNSWs*.

Accordingly, the query range $[L, R]$ is decomposed into a suffix range $[L, fs]$ contained in the left child segment and a prefix range $[fs+1, R]$ contained in the right child segment. This decomposition reduces arbitrary range search to two one-sided subrange searches and one parent segment for connectivity.

Example 1. Figure 1 illustrates the first-split range representation. The segment tree is built over the attribute domain $[1, 16]$, and the query range is $[7, 14]$. Since $[7, 14]$ crosses the midpoint $fs = 8$ of the root segment $[1, 16]$, the root is the *parent segment*, and $fs = 8$ is the *first-split position*.

Accordingly, the query range is split into two one-sided ranges: the suffix range $[7, 8]$ in the left child $[1, 8]$, and the prefix range $[9, 14]$ in the right child $[9, 16]$. Therefore, FRASP only needs to access the suffix HNSW restricted to $[7, 8]$, the prefix HNSW restricted to $[9, 14]$, and the parent HNSW on $[1, 16]$ restricted to $[7, 14]$. Other lower-level segments are skipped, even if they intersect the query range.

B. Suffix and Prefix HNSW Indexing

Note that dedicated techniques are unnecessary for very small ranges (e.g., containing fewer than $B = 1024$ objects), as such queries can be efficiently handled by pre-filtering. Thus, for the segment tree, we truncate levels with segment size $\leq B/2$, resulting in $O(n/B)$ nodes and height $O(\log(n/B))$. This suffices for all query ranges larger than B , since our design only requires searching the three segments at




Fig. 1. Suffix/Prefix structure and the first-split position of query [7, 14] in a segment tree.

the first-split level, and any such query must have a first-split segment size $> B/2$.

For each segment, we construct a specialized HNSW. If the segment is a left child of the segment tree, we construct a suffix HNSW for it. Otherwise, we construct a prefix HNSW for it. The root segment employs a prefix HNSW by default. Next, we detail the construction of the suffix HNSW for a left child segment $[l, r]$, and the construction of a prefix HNSW for a right child is analogous.

Suffix HNSW Construction. The suffix HNSW consists of (i) an HNSW graph index $\mathcal{H}_{\text{sfx}}$ and (ii) a pruned edge index N^{del} that stores the edges pruned during graph construction. Given a segment $[l, r]$ which is a left child in the segment tree, $\mathcal{H}_{\text{sfx}}$ is constructed by inserting objects in suffix order, i.e., from r down to l in descending order of attribute values. During the insertion of each object o_i , the pruned edges are recorded in N^{del} with a *deletion tag* $= i$ for each pruned edge, indicating that the edge was deleted during the insertion of o_i .

Algorithm 2 describes the construction of the suffix HNSW for $[l, r]$. After initializing empty $\mathcal{H}_{\text{sfx}}$ and N^{del} , we also create an array A to record the topmost layer where each object starts being inserted, which will be used in Section IV-C). Objects (e.g., o_i) are inserted in suffix order from r down to l (Line 3). In Line 4, similar to Lines 1-4 in Algorithm 1, we sample a random layer $\mathcal{L}_a$ and obtain the entry point ep_a . Downwards from $\mathcal{L}_a$, we start inserting o_i . The topmost layer $\mathcal{L}_a$ of inserting o_i is recorded in $A[i]$ at Line 5. For each layer $\mathcal{L}_b$ from $\mathcal{L}_a$ to $\mathcal{L}_0$, by Lines 6 to 10 we create new edges for o_i (i.e., $N_b[o_i]$), and prune existing edges of every neighbor o_j in $N_b[o_i]$ using the new edge (o_i, o_j) . The key point is that, before pruning, we copy the existing neighbors of o_j into $N'_b[o_j]$ (Line 9). The difference between $N'_b[o_j]$ and the updated neighbors $N_b[o_j]$ identifies the pruned edges (o_j, o_t) (Line 11), which are stored in $N_b^{\text{del}}[o_j]$ with a deletion tag i (Line 12). The setting $\text{tag} = i$ indicates that the edge (o_j, o_t) was deleted during the insertion of o_i into $\mathcal{H}_{\text{sfx}}$. If the edge (o_j, o_t) lies within the suffix subrange $[L, r]$ (i.e., o_j and o_t are both in $[L, r]$, and has $\text{tag} = i < L$), then the edge was deleted due to an insertion outside the query range. Thus, the edge remains valid for the range $[L, r]$.

Lemma 1 states that the in-range edges in $\mathcal{H}_{\text{sfx}}$, combined




Fig. 2. An example on inserting objects to a suffix HNSW.

with the in-range pruned edges in N^{del} where $\text{tag} = i < L$, collectively form an HNSW equivalent to one constructed directly over $[L, r]$ using the same M and ef_{con} parameters.

Lemma 1. *Given a suffix HNSW for segment $[l, r]$, consisting of $\mathcal{H}_{\text{sfx}}$ and N^{del} , for any suffix query range $[L, r]$ ($l \leq L \leq r$), the HNSW induced by objects in $[L, r]$ is equivalent to the HNSW that would be constructed directly over $[L, r]$ with the same insertion order, M , and ef_{con} .*

Example 2. Figure 2 illustrates how a suffix HNSW is maintained and reused for different suffix ranges. Figure 2(1) first depicts the suffix HNSW for range $[5, 8]$, where objects o_8, o_7, o_6, o_5 have been inserted in descending attribute order.

Figure 2(2) then shows the effect of inserting o_4 . The maintained range is extended to $[4, 8]$; a new edge (o_4, o_5) is added; and the existing edge (o_5, o_7) is pruned by the standard HNSW neighbor selection due to the degree constraint. FRASP stores this removed edge in N^{del} with deletion tag 4, indicating that the edge is removed when o_4 is inserted.

The two cases in Figures 2(3a) and 2(3b) show how the same maintained suffix HNSW serves different query ranges. For query range $[4, 8]$, o_4 is in range, so its insertion is valid and the current graph is used directly. For query range $[5, 8]$, o_4 is out of range; hence the pruning effect caused by o_4 should be undone. Since the stored tag on the pruned edge satisfies $4 < L = 5$, FRASP restores (o_5, o_7) . This allows FRASP to represent multiple suffix ranges with one maintained suffix HNSW and tagged pruned edges, rather than materializing an independent graph for each range.

Prefix HNSW Construction. The construction of the prefix HNSW for a right child segment $[l, r]$ is analogous to the suffix HNSW, except that objects are inserted in ascending order from l to r , as shown in Algorithm 3. In the prefix HNSW, a pruned edge (o_j, o_t) with $\text{tag} = i$ indicates it was deleted during the insertion of o_i . If an edge (o_j, o_t) lies within the prefix query range $[l, R]$ and has $\text{tag} = i > R$, then the edge was deleted due to an insertion outside the query range. Therefore, the edge remains valid for the query range $[l, R]$.

Algorithm 2: BuildSuffixHNSW($[l, r]$)

Input: Segment $[l, r]$, objects $\{o_l, \dots, o_r\}$, M , ef_{con}
Output: $\mathcal{H}_{\text{sfx}}$ and N^{del} for $[l, r]$ and A

- 1 Initialize empty $\mathcal{H}_{\text{sfx}}$ and N^{del} ;
- 2 Initialize A as a zero-array for indices in $[l, r]$;
- 3 **for** i from r down to l **do**
- 4 $\mathcal{L}_a, ep_a \leftarrow$ Same as Lines 1-4 in Algorithm 1;
- 5 $A[i] \leftarrow a$; // Record the topmost layer of o_i
- 6 **for** $\mathcal{L}_b$ from $\mathcal{L}_a$ to $\mathcal{L}_0$ **do**
- 7 $N_b[o_i] \leftarrow$ Same as Lines 6-7 in Algorithm 1;
- 8 **foreach** $o_j \in N_b[o_i]$ **do**
- 9 $N'_b[o_j] \leftarrow N_b[o_j]$;
- 10 $N_b[o_j] \leftarrow \text{PruneExistingEdges}(N_b[o_j] \cup \{o_i\})$;
- 11 **foreach** $o_t \in N'_b[o_j] \setminus N_b[o_j]$ **do**
- 12 $N_b^{\text{del}}[o_j].\text{add}((o_t, \text{tag} = i))$;
- 13 **return** $\mathcal{H}_{\text{sfx}}$, N^{del} , A ;

Algorithm 3: BuildPrefixHNSW($[l, r]$)

Input: Segment $[l, r]$, objects $\{o_l, \dots, o_r\}$, M , ef_{con}
Output: $\mathcal{H}_{\text{pfx}}$ and N^{del} for $[l, r]$

- 1 Initialize empty $\mathcal{H}_{\text{pfx}}$ and N^{del} ;
- 2 Initialize A as a zero-array for indices in $[l, r]$;
- 3 **for** i from l to r **do**
 // Same as Algorithm 2

The prefix HNSW satisfies a symmetric property to Lemma 1.

Discussion. While the segment graph in SeRF [10] shares the idea of fixing insertion orders to support one-sided query ranges, FRASP adopts different designs to handle arbitrary query ranges. Specifically, SeRF extends to 2DSegmentGraph for arbitrary query ranges by embedding four attribute-related tags per edge for bound checking. In contrast, FRASP only requires a single tag per pruned edge. Moreover, FRASP efficiently supports arbitrary ranges by enhancing a segment tree with prefix and suffix HNSWs. The performance improvements in experiments also validate the effectiveness of our design.

C. Succinct Array Indexes

Given a query range $[L, R]$, the top layers of an HNSW may contain no object in $[L, R]$ due to their sparsity. Thus, we must identify the highest layer containing at least one in-range object (i.e., the *topmost query-relevant layer* $\mathcal{L}_{\text{qmax}}$) and select such an object as the *query-relevant entry point* ep_{qmax} .

Recall that for any $[L, R]$ within a parent segment $[l, r]$, the midpoint mid is the first-split position fs . When building a specialized HNSW for $[l, r]$, we also recorded the highest layer where each object is inserted in an array A (see Line 5 in Algorithm 2). The midpoint mid splits A into two parts, $A[l : mid]$ and $A[mid + 1 : r]$, corresponding to the suffix and prefix subranges of the query range $[L, R]$. Leveraging this information, we build a succinct array index for each segment, to identify $\mathcal{L}_{\text{qmax}}$ and ep_{qmax} for $[L, R]$ in the parent HNSW through simple array lookups during search.

Algorithm 4: BuildSuccinctArrayIndex

Input: Array A for parent segment $[l, r]$
Output: suffix-Max array Z_1 , prefix-Max array Z_2

- 1 $mid \leftarrow \lfloor (l + r) / 2 \rfloor$;
- 2 Initialize Z_1 for $[l, mid]$, Z_2 for $[mid + 1, r]$ as zero-arrays;
 // Build suffix-Max array
- 3 $Z_1[mid] \leftarrow mid$;
- 4 **for** i from $mid - 1$ down to l **do**
- 5 **if** $A[i] > A[Z_1[i + 1]]$ **then**
- 6 $Z_1[i] \leftarrow i$;
- 7 **else**
- 8 $Z_1[i] \leftarrow Z_1[i + 1]$;
- 9 // Build prefix-Max array
- 10 $Z_2[mid + 1] \leftarrow mid + 1$;
- 11 **for** i from $mid + 2$ to r **do**
- 12 **if** $A[i] > A[Z_2[i - 1]]$ **then**
- 13 $Z_2[i] \leftarrow i$;
- 14 **else**
- 15 $Z_2[i] \leftarrow Z_2[i - 1]$;
- 16 **return** Z_1, Z_2 ;

Specifically, with $A[l : mid]$, we construct a *suffix-Max array* Z_1 by a linear scan from mid down to l in A . Here, $Z_1[i]$ stores the index of the object with the highest layer among all objects in the suffix range $[i, mid]$, where i is an index in $[l, mid]$. Thus, the topmost query-relevant layer for the suffix query range $[L, mid]$ is $A[Z_1[L]]$, and the query-relevant entry point is $Z_1[L]$. Similarly, for $A[mid + 1 : r]$, we build a *prefix-Max array* Z_2 by a linear scan from $mid + 1$ up to r in A . Here, $Z_2[i]$ stores the index of the object with the highest layer among all objects in the prefix range $[mid + 1, i]$, where i is an index in $[mid + 1, r]$. Thus, the topmost query-relevant layer for the prefix query range $[mid + 1, R]$ is $A[Z_2[R]]$, and the query-relevant entry point is $Z_2[R]$.

With suffix-Max and prefix-Max arrays, we determine $\mathcal{L}_{\text{qmax}}$ and ep_{qmax} by Eq. (1). That is, during query processing, we compare the entry points identified by the suffix-Max and prefix-Max arrays, and select the one with the higher layer as the final $\mathcal{L}_{\text{qmax}}$ and ep_{qmax} for $[L, R]$ in the parent HNSW.

$$\begin{aligned} \mathcal{L}_{\text{qmax}} &= \max(A[Z_1[L]], A[Z_2[R]]) \\ ep_{\text{qmax}} &= \begin{cases} Z_1[L], & \text{if } A[Z_1[L]] > A[Z_2[R]] \\ Z_2[R], & \text{otherwise} \end{cases} \end{aligned} \quad (1)$$

Algorithm 4 presents the pseudo code of constructing the suffix-Max and prefix-Max arrays Z_1 and Z_2 for the parent segment $[l, r]$ from A . After computing mid and initializing the two arrays in Lines 1–2, Lines 3–8 build Z_1 by scanning from mid down to l . The base case is $Z_1[mid] = mid$. For each $i \in [l, mid - 1]$ scanned in descending order, we compare $A[i]$ with $A[Z_1[i + 1]]$. If $A[i]$ is larger, we set $Z_1[i] = i$; otherwise, we inherit the previous argmax, i.e., $Z_1[i] = Z_1[i + 1]$. Hence, $Z_1[i]$ stores the index of the highest-layer object in the suffix range $[i, mid]$.

Lines 9–14 build Z_2 symmetrically by scanning from $mid + 1$ up to r . The base case is $Z_2[mid + 1] = mid + 1$. For each $i \in [mid + 2, r]$, we compare $A[i]$ with $A[Z_2[i - 1]]$. If $A[i]$ is larger, we set $Z_2[i] = i$; otherwise, we set $Z_2[i] = Z_2[i - 1]$.




Fig. 3. Find query-relevant entry point through the succinct arrays.

Thus, $Z_2[i]$ stores the index of the highest-layer object in the prefix range $[mid + 1, i]$.

Example 3. Figure 3 illustrates how FRASP locates an in-range entry point in the parent HNSW. For query range $[3, 7]$ in parent segment $[1, 8]$, the midpoint $mid = 4$ splits it into $[3, 4]$ and $[5, 7]$. Z_1 and Z_2 shown in the figure correspond to the suffix-Max array and the prefix-Max array, respectively.

By array lookup, $Z_1[3] = 3$ gives object o_3 with layer $A[3] = 1$, while $Z_2[7] = 5$ gives object o_5 with layer $A[5] = 2$. Since $A[5] > A[3]$, Eq. (1) sets $ep_{qmax} = o_5$ and $\mathcal{L}_{qmax} = \mathcal{L}_2$. Thus, FRASP starts the upper-layer search from the highest in-range object instead of the general HNSW entry point.

D. Overall Indexing and Analysis

The overall procedure of FRASP constructing the index is provided in Algorithm 5. We first construct a segment tree over the input data range $[1, n]$ (Line 1), which is a standard procedure and thus the details are omitted. For each segment in this tree, if it is a left child, we build a suffix HNSW $\mathcal{H}_{sfx}$ with N^{del} pruned edges (Lines 3-4). If the segment is a right child or the root, it constructs a prefix HNSW $\mathcal{H}_{pfx}$ with N^{del} pruned edges (Lines 5-6). For non-leaf segments that can serve as parents during query processing, we build suffix-max and prefix-max arrays Z_1 and Z_2 (Lines 7-8).

Construction Time of Prefix/Suffix HNSW. The time complexity of constructing a prefix or suffix HNSW is equivalent to that of a standard HNSW. For a standard HNSW with n objects and parameter M , the total number of edges is $O(nM)$. During construction, the procedure of pruning edges may delete up to M edges, leading to a total of $O(nM)$ pruned edges for a single HNSW in the worst case. Thus, the overall space complexity of a prefix or suffix HNSW, including pruned edges, remains $O(nM)$, consistent with the original HNSW. Empirically, the number of pruned edges is often a small fraction of the total edges, as shown in the experiments.

Space Complexity. For a standard HNSW with n objects and parameter M , the number of edges is $O(nM)$. More precisely, the expected number of edges is $O(nM(1 + 1/\ln(M)))$, where the term $M(1 + 1/\ln(M))$ accounts for the expected number of edges per object across all layers in an HNSW [7]. Usually M is 16 to 64, meaning that $1/\ln(M) < 1$, and thus it

Algorithm 5: FRASP Index Construction

Input: Objects $\mathcal{D}$, HNSW parameters M, ef_{con}, B

Output: FRASP-Index

```

1 Build a segment tree over  $[1, n]$  with  $B$ ;
2 foreach segment  $[l, r]$  in the segment tree do
3   if  $[l, r]$  is a left child then
4      $\mathcal{H}_{sfx}, N^{\text{del}}, A \leftarrow \text{BuildSuffixHNSW}([l, r], M, ef_{\text{con}})$ ;
5   else if  $[l, r]$  is a right child or the root segment then
6      $\mathcal{H}_{pfx}, N^{\text{del}}, A \leftarrow \text{BuildPrefixHNSW}([l, r], M, ef_{\text{con}})$ ;
7   if  $[l, r]$  is not a leaf then
8     // It can serve as a parent segment
8      $Z_1, Z_2 \leftarrow \text{BuildSuccinctArrayIndex}(A)$ ;

```

is reduced to $O(nM)$. Then, since the objects in different segments of the same level in the segment tree are disjoint, the total edge complexity of HNSWs in the same level is $O(nM)$. The total number of pruned edges for each level of the segment tree is also $O(nM)$ in the worst case. As there are $O(\log(n/B))$ levels in the segment tree, the total number of edges in the whole FRASP index is $O(nM \log(n/B))$, reduced to $O(nM \log n)$ if B is a constant. For every segment, we build arrays A , Z_1 , and Z_2 , all of which together account for $O(n \log n)$ space. Therefore, the total space complexity of FRASP is $O(nM \log n)$.

Runtime Analysis. The main computational cost of HNSW indexing comes from searching through its layers. Although it has been shown that no algorithm can guarantee sub-linear time complexity for nearest neighbor search [42]–[44], graph-based indexing methods often exhibit sub-linear search performance in practice on real-world datasets [9]. Therefore, theoretical asymptotic complexity may not fully capture practical performance. Following the assumption in [9], we focus on analyzing the additional overhead introduced by our index construction compared to building a single HNSW over the entire object set $\mathcal{D}$.

Let $h(n)$ denote the time complexity to build a single HNSW, and $f(n)$ denote the time complexity to search during edge creation, where n is the number of objects. By [9], we have $f(n) = O(n^c)$ where $\frac{1}{2} < c \leq 1$. Note that, FRASP builds 2^i disjoint HNSWs of size $n/2^i$ for $t = \log(n/B) + 1$ levels of the segment tree. Then, we derive the time complexity of generating the indexes by FRASP in Theorem 4.

Theorem 4 (Indexing Time Complexity). *The time complexity of FRASP constructing indexes is $h(n) \log(n/B)$. For the search process during edge creation, the runtime of FRASP is bounded by $f(n) \cdot \frac{1 - (1/2^c)^{\log(n/B)+1}}{1 - 1/2^c}$, whose maximum is $\frac{1}{1 - \sqrt{2}/2}$ obtained at $c = \frac{1}{2}$.*

V. QUERY PROCESSING

With the indexes built, we are ready to detail the processing of queries. Given a query vector q and range $[L, R]$, we first locate the first-split position $fs = mid = \lfloor (l + r)/2 \rfloor$ in the segment tree, where $[l, r]$ is the minimal segment containing

Algorithm 6: FRASP Query Processing

Input: Query vector q , query range $[L, R]$, number of neighbors k , candidate size ef

Output: k nearest neighbors to q within range $[L, R]$

```

1  $fs = mid, l, r \leftarrow \text{FindFirstSplitPosition}([L, R]);$ 
2  $\mathcal{H}_{\text{sfx}}, \mathcal{H}_{\text{pfx}}, \mathcal{H}_p \leftarrow \text{HNSWs for } [l, mid], [mid + 1, r], [l, r];$ 
3  $\mathcal{L}_{\text{qmax}}, ep_{\text{qmax}} \leftarrow \text{Eq. (1) on } Z_1, Z_2, A \text{ of } \mathcal{H}_p;$ 
4  $ep_p \leftarrow \text{FindBaseEP}(\mathcal{H}_p, q, \mathcal{L}_{\text{qmax}}, ep_{\text{qmax}});$ 
5 Initialize priority queues  $\mathcal{C}$  and  $\mathcal{R}$  with  $(ep_p, \delta(q, ep_p));$ 
  // Shared search process on base layers
6 while  $\mathcal{C}$  is not empty do
7    $o_c \leftarrow \mathcal{C}.\text{pop}()$  // Pop the closest candidate
8   if  $\delta(q, o_c) > \mathcal{R}.\text{maxDist}$  and  $|\mathcal{R}| \geq ef$  then break;
  // Find the neighbors of  $o_c$  to be explored
9   Set  $\mathcal{V} \leftarrow \emptyset;$ 
10  foreach  $\mathcal{H}$  of  $\mathcal{H}_{\text{pfx}}, \mathcal{H}_{\text{sfx}},$  and  $\mathcal{H}_p$  where  $o_c$  exists do
11    foreach neighbor  $o_j \in \mathcal{H}.N_0[o_c]$  do
12      if  $L \leq j \leq R$  and  $o_j$  is unvisited then
13         $\mathcal{V}.\text{insert}(o_j);$ 
14        Mark  $o_j$  as visited;
15      foreach neighbor  $o_j \in \mathcal{H}.N_0^{\text{del}}[o_c]$  do
16        if  $((\mathcal{H} \text{ is suffix HNSW and } (o_i, o_j).\text{tag} < L) \text{ or}$ 
17           $(\mathcal{H} \text{ is prefix HNSW and } (o_i, o_j).\text{tag} > R))$  and
18           $L \leq j \leq R$  and  $o_j$  is unvisited then
19             $\mathcal{V}.\text{insert}(o_j);$ 
20            Mark  $o_j$  as visited;
21      foreach neighbor  $o_j \in \mathcal{V}$  do
22        if  $\delta(q, o_j) < \mathcal{R}.\text{maxDist}$  or  $\mathcal{R}.\text{size} < ef$  then
23           $\mathcal{C}.\text{insert}(o_j, \delta(q, o_j));$ 
24           $\mathcal{R}.\text{insert}(o_j, \delta(q, o_j));$ 
25          if  $\mathcal{R}.\text{size} > ef$  then  $\mathcal{R}.\text{pop}();$ 
26 return Top- $k$  in  $\mathcal{R};$ 

```

$[L, R]$ (also, the parent segment). The query range is then split into the suffix range $[L, mid]$ and prefix range $[mid + 1, R]$. The suffix HNSW $\mathcal{H}_{\text{sfx}}$ supports $[L, mid]$; the prefix HNSW $\mathcal{H}_{\text{pfx}}$ supports $[mid + 1, R]$; and the parent HNSW $\mathcal{H}_p$ over $[l, r]$ provides inter-segment connectivity. In a special case where the query range $[L, R]$ has $R = mid$ (resp. $L = mid + 1$), we only need to search one HNSW, the suffix HNSW $\mathcal{H}_{\text{sfx}}$ (resp. prefix HNSW $\mathcal{H}_{\text{pfx}}$). Next, we only focus on the general case where $L \leq mid$ and $R > mid$.

For efficient query processing, FRASP employs three ideas. (i) A *single unified search process* is performed over the base layers of $\mathcal{H}_{\text{sfx}}, \mathcal{H}_{\text{pfx}},$ and $\mathcal{H}_p$, using a shared candidate priority queue $\mathcal{C}$ and result priority queue $\mathcal{R}$, with a global lower bound $\mathcal{R}.\text{maxDist}$ for pruning and early termination. (ii) The search is strictly *in-range*, starting from entry points found via the suffix-max and prefix-max arrays in Section IV-C, and explores only objects within $[L, R]$, ensuring efficiency even for small overlaps with the segments. (iii) Any edge between objects o_i and o_j may appear in both $\mathcal{H}_p$ and either $\mathcal{H}_{\text{sfx}}$ or $\mathcal{H}_{\text{pfx}}$, but is *traversed at most once* during the search, avoiding redundant computations.

The detailed process to perform the query by FRASP is presented in Algorithm 6. We first find the first-split position $fs = mid$ of query range $[L, R]$ and the parent segment $[l, r]$, by applying a routine called `FindFirstSplitPosition`,

which utilizes a simple binary search in the segment tree. Thus, its pseudocode is omitted here but deferred to Algorithm 8 in our technical report [17]. Next, we derive the left child $[l, mid]$, right child $[mid + 1, r]$, and parent $[l, r]$ segments associated with their HNSWs ($\mathcal{H}_{\text{sfx}}, \mathcal{H}_{\text{pfx}},$ and $\mathcal{H}_p$, resp.) in Line 2. In Lines 3-4, for the parent HNSW $\mathcal{H}_p$, we apply Eq. (1) to determine its topmost query-relevant layer and the query-relevant entry point, which are then passed to a routine `FindBaseEP` in Line 4, which locates an in-range nearest neighbor in $\mathcal{L}_1$ as the base-layer entry point ep_p . In Line 5, we initialize a min-priority queue $\mathcal{C}$ to store possible nearest neighbors (called “candidates”) and a max-priority queue $\mathcal{R}$ to store the actual nearest neighbors, by initializing them with the entry point of the base-layer ep_p ², along with their distances to the query q .

The details of `FindBaseEP` are presented in Algorithm 7. `FindBaseEP` performs a search on $\mathcal{H}_p$ from $\mathcal{L}_{\text{qmax}}$ down to $\mathcal{L}_1$, updating ep to the closest in-range neighbor per layer using a flag *changed* to track improvements. It only considers in-range neighbors and valid pruned edges. The process terminates when no closer neighbor is found, returning the final ep as the base-layer entry point. `FindBaseEP` is efficient, incurring negligible overhead compared to the main base-layer search, as it only traverses the middle layers to locate a single in-range entry point.

Next, we conduct a single unified search process over the base layers of $\mathcal{H}_{\text{sfx}}, \mathcal{H}_{\text{pfx}},$ and $\mathcal{H}_p$, by iterating over the candidates in $\mathcal{C}$. Each time, we take out the candidate closest to q , denoted as o_c (Line 7). If the distance of the closest candidate to q is larger than $\mathcal{R}.\text{maxDist}$ and $|\mathcal{R}| \geq ef$, we simply stop iterating through $\mathcal{C}$ (Line 8), just like HNSW [7]. Otherwise, we proceed and denote the HNSW among $\mathcal{H}_{\text{pfx}}, \mathcal{H}_{\text{sfx}},$ and $\mathcal{H}_p$ which contains o_c as $\mathcal{H}$.

In Lines 11–18, we try to insert the neighbors of o_c in $\mathcal{H}$ into $\mathcal{V}$. Specifically, for each neighbor o_j of o_c (i.e., $o_j \in N_0[o_c]$ in the base layer of $\mathcal{H}$), if o_j is in-range (i.e., $L \leq o_j.a \leq R$) and unvisited, we insert it into $\mathcal{V}$ and mark it as visited (Lines 11–14). To recover the deleted neighbors that should have been in the query range, for each deleted neighbor o_j in $N_0^{\text{del}}[o_c]$, if o_j is in-range, unvisited, and its deletion was caused by inserting an object outside the query range (i.e., $o_j.\text{tag} < L$ when $\mathcal{H}$ is a suffix HNSW or $o_j.\text{tag} > R$ when $\mathcal{H}$ is a prefix HNSW), we insert it into $\mathcal{V}$ and mark it as visited.

After obtaining the set $\mathcal{V}$, we iterate through each object o_j in $\mathcal{V}$ to update $\mathcal{C}$ and $\mathcal{R}$. If the distance from q to o_j (i.e., $\delta(q, o_j)$) is smaller than the maximum distance in $\mathcal{R}$, we insert o_j into both $\mathcal{C}$ and $\mathcal{R}$. After the insertion, if the size of $\mathcal{R}$ exceeds ef , we pop an object from $\mathcal{R}$ to keep its size no larger than ef (Lines 21–23).

The above iteration process goes on, until the candidate queue is empty. When the iteration stops, we return the top- k nearest neighbors in $\mathcal{R}$ (Line 24).

²In the special case where the query range is exactly one segment, only one ep_{sfx} or ep_{pfx} is used instead.

Algorithm 7: FindBaseEP

Input: HNSW $\mathcal{H}$, q , $[L, R]$, $\mathcal{L}_{q_{\max}}$, $ep_{q_{\max}}$
Output: base-layer entry point ep

```

1  $ep \leftarrow ep_{q_{\max}}$ ;
2 for  $\mathcal{L}_b$  from  $\mathcal{L}_{q_{\max}}$  down to  $\mathcal{L}_1$  do
3    $changed \leftarrow \text{true}$ ;
4   while  $changed$  do
5      $changed \leftarrow \text{false}$ ;
6     foreach neighbor  $o_j \in \mathcal{H}.N_b[ep]$  do
7       if  $L \leq j \leq R$  and  $\delta(q, o_j) < \delta(q, ep)$  then
8          $ep \leftarrow o_j$ ;  $changed \leftarrow \text{true}$ ;
9     foreach neighbor  $o_j \in \mathcal{H}.N_b^{del}[ep]$  do
10      if  $L \leq j \leq R$  and (( $\mathcal{H}$  is suffix HNSW and
        ( $o_i, o_j$ ).tag <  $L$ ) or ( $\mathcal{H}$  is prefix HNSW and
        ( $o_i, o_j$ ).tag >  $R$ )) then
11        if  $\delta(q, o_j) < \delta(q, ep)$  then
12           $ep \leftarrow o_j$ ;  $changed \leftarrow \text{true}$ ;
13 return entry point  $ep$  for base layer;

```

Time Complexity. The asymptotic query complexity is not a very informative measure for comparing graph-based methods, since their worst-case runtimes are typically linear to n , while their practical performance differs significantly. Note that their practical performance is mainly determined by the number of objects processed during search. Therefore, just like [9], we analyze the time to process per object. Theorem 5 shows that FRASP requires an expected $O(M)$ time per processed object.

Theorem 5. *The expected time complexity to process the edges of an object in FRASP is $O(M)$.*

Adaptation to Dynamic Datasets. The current design of FRASP is primarily optimized for static or batch-updated datasets. In real-world applications, vector databases often face dynamic workloads requiring frequent insertions and deletions [45]. Thus, we would like to extend FRASP to dynamic scenarios in our future work. Here, we provide the main ideas of extending FRASP. (1) To cope with the deletion of objects, we will employ a soft-delete (tombstone) mechanism. By maintaining a global bitset, vectors can be marked as deleted in $O(1)$ time and filtered out during the query phase. A background thread will rebuild the affected sub-trees once the deletion ratio exceeds a specific threshold. (2) For random insertions, we will introduce a lightweight memory buffer (inspired by LSM-trees) for the incoming vectors [46]. Queries will search both FRASP and the buffer simultaneously, and a background process will merge the buffer into the main segment tree when it reaches capacity. In a special case where the attribute of the object to be inserted exceeds all existing objects, to be more efficient, we can simply allocate new leaf nodes for the incoming data and dynamically link them to the root, which completely avoids modifying the historical structures.

Extension to Multiple Attributes. Besides dynamic updates, accommodating queries with multiple attribute constraints is a natural extension with real-world applications. For example, in e-commerce, customers would like to select products by

TABLE II
MEMORY FOOTPRINT

	WIT	Sift	Gist	YT-rgb	Deep10M
Dimension	2048	128	960	1024	96
Raw Vectors	7.7	0.47	3.58	3.9	3.62
FRASP	9.13	2.24	5.01	4.50	26.23
iRangeGraph	12.87	1.97	8.81	9.05	26.80
SuperPostfiltering	28.81	7.34	16.64	17.36	91.90
SeRF	7.94	0.81	3.87	4.18	7.70
Milvus	8.41	0.78	4.05	4.31	7.22
ACORN	9.39	2.23	5.33	5.57	14.20
UNIFY	16.33	2.03	8.23	8.71	17.91
WoW	8.41	1.25	4.36	4.59	12.56

TABLE III
NUMBER OF PRUNED EDGES AND TOTAL EDGES IN THE INDEX.

Dataset	WIT	Sift	Gist	YT-rgb	Deep10M
Pruned Edges ($\times 10^6$)	33.20	49.15	30.10	2.23	926.73
Total Edges ($\times 10^6$)	176.50	216.18	167.60	57.30	3471.30
Ratio (%)	18.8	22.7	18.0	3.8	26.70

both timestamp and category. To handle such complex filters, we extent FRASP according to the type of attributes. (1) For queries involving multiple numerical attributes, FRASP will continue to build the segment tree based on the most selective continuous attribute. (2) For categorical attributes, we will maintain compressed bitsets or hard range boundaries. During the HNSW graph traversal, candidates will be efficiently validated against these secondary conditions via $O(1)$ bitwise AND operations. (3) For queries mixed with numerical and categorical attributes, we can first process the more selective attributes with FRASP, and then perform post-filtering for the other attributes.

VI. EXPERIMENTS

All methods are implemented in C++ and compiled with GCC 11.4.0 using `-O3` and `-march=native`. Experiments are conducted on an Intel Xeon Platinum 8370C (2.80GHz, 384GB RAM) running Ubuntu 22.04. All search performance results use a single thread. Code is available at <https://github.com/wjy-yy/FRASP>.

A. Experimental Settings

Datasets. The algorithms are evaluated on five widely adopted benchmarking vector datasets [9], [10], [41], [47]: WIT [48] (2048-dim ResNet50 image embeddings), YT-rgb [49] (1024-dim YouTube8M video features [50]), Sift [51] (128-dim image features), Gist [51] (960-dim image features), and Deep10M [52] (96-dim deep descriptors from the GoogleNet). Attributes are: image size for WIT, number of likes for YT-rgb [9], random integers in $[1, 10^4]$ for Sift and Gist, id for Deep10M [52]. Each dataset contains one million objects except Deep10M, which contains 10^7 objects. One thousand vectors are sampled as queries, following [9].

Query Workloads. We evaluate different query workloads based on the range filter size, defined as the fraction of

TABLE IV
TIME BREAKDOWN OF ALL QUERIES IN SECONDS.

Dataset	WIT	Sift	Gist	YT-rgb	Deep10M
Entry points	0.032	0.006	0.020	0.028	0.009
Dist computation	5.391	1.221	3.919	6.790	1.915
Graph traversal	1.249	1.028	1.332	1.658	1.600
Queue management	0.311	0.278	0.330	0.352	0.313
Total query time	6.984	2.534	5.601	8.828	3.837

data objects covered by the query range [9]. Specifically, a range fraction of 2^{-i} means the query range covers $n/2^i$ ($i \in \{2, 4, 6, 8\}$) data objects. In the experiments, we first consider a mixed workload where the query vectors are randomly partitioned into 10 subsets. For the i -th subset ($i \in [0, 9]$), each query vector is assigned a random range with fraction 2^{-i} . These 10 subsets together form the mixed workload. Second, we consider fixed workloads, where all query vectors share the same range fraction. By varying i , we evaluate performance under large, medium, and small query ranges.

Evaluation Metrics. We are concerned about two key metrics: the efficiency and the accuracy. Specifically, the efficiency of algorithms is measured by *queries per second (QPS)*, i.e., the number of queries processed per second [9], [47]. The accuracy of algorithms is measured by three metrics: *recall*, *distance relative error (DRE)*, and *total distance ratio (TDR)* [53]. Let Φ^* be the ground-truth k nearest neighbors and Φ the approximate results. Then, the recall is $|\Phi \cap \Phi^*|/|\Phi^*|$. DRE computed as $\frac{1}{k} \sum_{i=1}^k [\delta(o_i, q) - \delta(o_i^*, q)]/\delta(o_i^*, q)$, where o_i is the i -th ANN and o_i^* the i -th NN. TDR is given by $\sum_{i=1}^k \delta(o_i, q) / \sum_{i=1}^k \delta(o_i^*, q)$.

Methods and Settings. We compare FRASP with iRangeGraph [9], WoW [14], UNIFY [15], SuperPostfiltering [41], SeRF [10], ACORN [13], Milvus [11], and Pre-filtering, as reviewed in Section III. For all methods, $ef_{con} = 400$ unless otherwise specified. For all graph-based methods, we vary ef to adjust the QPS-Recall trade-off during query processing.

For FRASP, we set $M = 64$ for YT-rgb and $M = 16$ for the rest datasets, $B = 2^{10}$ for the first four datasets and $B = 2^{14}$ for Deep10M. M is also varied to study its impact on FRASP.

For the baselines, we use their recommended or default parameters whenever available. Notably, Pre-filtering does not need parameter setting. iRangeGraph only recommended setting $m = 64$ for WIT and YT-rgb. Then, for Sift and Gist, we use grid search and set $M = 16$, and we also set $M = 16$ for Deep10M following [14]. For SuperPostfiltering and UNIFY, we adopt their suggested value of $ef_{con} = 500$.

B. Experimental Results

To synthetically demonstrate the performance of the algorithms, we report the results of the aforementioned metrics in pairs, specifically QPS vs. recall and QPS vs. DRE/TDR.

QPS vs. Recall. Figure 4 shows the QPS-Recall curves of all methods across all datasets. We observe that FRASP consistently achieves higher QPS and recall than all baselines across all datasets and query workloads, often by a substantial margin. Under the RFANNS setting, the performance

differences (e.g., Figure 4 with range fraction 2^{-2}) among iRangeGraph, WoW, SeRF and SuperPostfiltering are already considered as significant, but FRASP delivers even larger improvements over them. For the mixed workload over Sift, at recall 0.9, FRASP achieves about 19,700 QPS— $1.6\times$ and $4.6\times$ that of iRangeGraph and SeRF (12,200 and 4,300 QPS, respectively), $23\times$ and $130\times$ that of ACORN (840 QPS) and Milvus (150 QPS). SuperPostfiltering usually has relatively good performance but FRASP still outperforms it by a large margin. Pre-filtering always achieves recall 1.0 but with low QPS, often appearing as an isolated point in the plots. These results demonstrate that FRASP consistently delivers high recall and efficiency for RFANNS queries, validating the superiority of our index and query processing approach.

QPS vs. DRE/TDR. The results of QPS vs. DRE under the mixed workload are presented in Figure 5. Observe that FRASP consistently achieves the best QPS and relative error over existing methods, often with a notable margin. FRASP also consistently achieves a better QPS-TDR trade-off, with total distance ratios closer to the exact top- k results while maintaining higher throughput. This validates that FRASP can achieve low error of distances, via the technical designs proposed in Sections IV and V. The results under other workloads and QPS vs. TDR are deferred to our technical report [17] due to space limitations.

Memory Footprint. We report the memory footprint for all methods except Pre-filtering (which does not require indexing) in Table II, including both the index and raw vectors. Among the top-performing methods, FRASP achieves a comparable or lower memory footprint, validating the space complexity $O(nM \log n)$ in Section IV-D. Further, we only consider the indexes near the first-split position of a query in the segment tree. This allows us to truncate the segment tree, considering the queries with ranges larger than B of interest. In contrast, iRangeGraph requires a complete segment tree, as it may need to access small segments for larger range queries. Thus, our index memory footprint is smaller than iRangeGraph on most datasets. Table III shows the number of pruned edges and total edges (including all edges in N and all pruned edges stored in N^{del} , across all layers) in the index. Observe that the ratio of pruned edges to total edges is small, showing that storing the pruned edges is not a bottleneck in the index.

Time Breakdown. As explained in Algorithm 6, we first find the in-range base layer entry points with minimal cost, and then perform the unified search. For all queries per dataset, we report the total query time, entry point finding time, and distance computation time in Table IV under $ef = 1000$ and mixed workload. Observe that the time to find base layer entry points is indeed minimal. While distance computation dominates in most datasets (e.g., WIT), in some datasets (e.g., Sift), the overhead of other costs, including graph traversal and queue management, is also significant.

Comparison with Oracle-HNSW. Following [9], we compare FRASP with Oracle-HNSW, which builds a separate HNSW index for each query range—impractical due to its $O(n^3M)$




Fig. 4. QPS-Recall on WIT, Sift, Gist, YT-rgb, and Deep10M (top to bottom; mixed and range-fraction workloads left to right).




Fig. 5. QPS-Distance Relative Error on the Mixed Workload.




Fig. 6. Distance computations per query vs. recall on the mixed workload (y -axis in thousands; $\times 10^3$ in each panel).




Fig. 7. Vary M on the mixed workload at $EF=1000$: QPS, recall, distance relative error, and distance computations per query.




Fig. 8. Ablation Study.




Fig. 9. FRASP vs. Oracle-HNSW under mixed workloads.

space cost to support all possible $O(n^2)$ query ranges. We construct Oracle-HNSW for 100 randomly sampled queries per dataset, as indexing all queries is infeasible. Figure 9 shows QPS-Recall on WIT and Gist (mixed workload): FRASP achieves performance comparable to or even surpassing Oracle-HNSW while using only $O(nM \log n)$ space, demonstrating the efficiency of our approach.

Effect of M . We vary the maximum out-degree $M \in \{4, 8, 16, 32, 64\}$ on the mixed workload at $EF=1000$ and report QPS, recall, distance relative error, and per-query distance computations in Figure 7. Recall remains high across these settings. QPS, distance relative error and per-query distance computations vary substantially with M . Different datasets show different sensitivities, so the best M can differ by dataset.

Ablation Study. We conduct an ablation study to compare FRASP that uses three segments and their HNSWs, to Child-only (considering only the child segments) and Parent-only (considering only the parent segment). In *Child-only*, the two child HNSWs are disconnected, so a single search entry can traverse only one side; we therefore invoke the same entry-point selection routine *in each child* to obtain a starting vertex per child and begin the layered search from both. As shown in Figure 8 on mixed workload, the performance of Child-only is degraded, indicating that the disconnected HNSWs in the children are not sufficient for high recall. It requires the HNSW of parent segment to introduce inter-edges. The performance of Parent-only also drops because the parent HNSW is built as a suffix or prefix HNSW over a range $[l, r]$ that strictly

contains $[L, R]$. The ablation study shows the effectiveness of our design to use three segments and their HNSWs to achieve high QPS and recall over existing methods as reported above.

VII. CONCLUSION

In this paper, we propose FRASP, an efficient and effective method for RFANNS. We leverage a first-split range strategy combined with suffix and prefix HNSW indexes to enable fast and accurate RFANNS. By consulting only three segment-specific graph indexes, we develop a unified search process that efficiently restricts exploration to in-range objects without redundant computations, reducing the computation cost per object. We also design succinct prefix/suffix array indexes to quickly locate relevant HNSW layers and entry points. Our approach effectively balances the trade-off between query efficiency and result quality, and demonstrates excellent performance across various datasets and query workloads. As future work, we plan to extend our methods to support more complex range-filtering queries, such as queries on dynamic vector data.

VIII. PROOFS

A. Proof of Lemma 1

Proof. For this proof, we demonstrate that the combination of in-range edges from $\mathcal{H}_{\text{sfx}}$ and the relevant deleted edges from N^{del} with tag $i < L$ forms a graph structure equivalent to an HNSW built directly over objects in range $[L, r]$.

Let $\mathcal{H}_{\text{sfx}}^L$ represent the HNSW that would be built by inserting objects in descending order from r down to L (i.e., an HNSW built directly on the range $[L, r]$).

First, observe that due to the suffix order insertion (from r down to L), at any step i where $l \leq i < L$, the existing HNSW contains all objects $\{o_{i+1}, \dots, o_r\}$. When object o_i is inserted, it may cause the deletion of edges between objects in the range $[i+1, r]$. These deletions are recorded in N^{del} with tag i .

Consider any edge (o_j, o_k) where $L \leq j, k \leq r$ (i.e. both endpoints are in the range $[L, r]$). We analyze three cases for the edge (o_j, o_k) :

1. If edge (o_j, o_k) exists in the final $\mathcal{H}_{\text{sfX}}$, it would also exist in $\mathcal{H}_{\text{sfX}}^L$ since no insertion after o_L would affect it differently in the two construction processes.

2. If edge (o_j, o_k) was deleted during the insertion of o_i with $i < L$, this edge would exist in $\mathcal{H}_{\text{sfX}}^L$ because object o_i is not included in the construction of $\mathcal{H}_{\text{sfX}}^L$. Such edges are precisely those in N^{del} with tag $i < L$.

3. If edge (o_j, o_k) was deleted during the insertion of o_i with $i \geq L$, this deletion would occur in both $\mathcal{H}_{\text{sfX}}$ and $\mathcal{H}_{\text{sfX}}^L$ following the same pruning criteria, so the edge correctly does not appear in our reconstructed graph.

Therefore, by combining the in-range edges from $\mathcal{H}_{\text{sfX}}$ and the deleted edges from N^{del} with tag $i < L$, we precisely reconstruct the edge set of $\mathcal{H}_{\text{sfX}}^L$. The resulting graph maintains the same hierarchical structure and neighbor relationships as if it had been built directly on the objects in $[L, r]$ using the same parameters M and ef_{con} . $\square$

B. Lemma for Prefix HNSW

Lemma 2. Given a prefix HNSW $\mathcal{H}_{\text{pfx}}$ for segment $[l, r]$ and its pruned edges N^{del} , for any prefix query range $[l, R]$ ($l \leq R \leq r$), the HNSW induced by objects in $[l, R]$ —including all in-range edges of $\mathcal{H}_{\text{pfx}}$ and all in-range pruned edges in N^{del} with tag $i > R$ —is equivalent to the HNSW that would be constructed directly over $[l, R]$ with the same insertion order, M , and ef_{con} .

Proof of Lemma 2 is similar to that of Lemma 1.

C. Proof of Theorem 4

Proof. For the first part of the theorem, define $h(n)$ as the time complexity of single HNSW of size n . First, we have $\forall x, y, h(x) + h(y) \leq h(x + y)$. The left side is considered to establish HNSWs with x and y points separately, while the right side is to establish an HNSW with $x + y$ points. Since the right side is a single HNSW with more points, it may take longer to perform SearchLayer operations to find the nearest neighbors to create the edges. And every other time step is almost the same.

Assuming n is dividable by B , and 2^i for brevity to ignore rounding up and down, we can get that our segment tree has $t = \log(n/B) + 1$ layer, supposing from top to down starting from index i . In layer i , we have 2^i elements whose size is $\frac{n}{2^i}$. Then FRASP index time is:

$$\sum_{i=0}^{t-1} (2^i) h\left(\frac{n}{2^i}\right) \leq \sum_{i=0}^{t-1} h(n) = h(n) \log(n/B)$$

For the second part, following the assumption of made in [9], let $f(n) = O(n^c)$ be the time complexity of SearchLayer on a n objects NSW graph based on its performance on the real-world dataset. The expected number of objects on

$\mathcal{L}_a$ layer in HNSW is (according to the sample probability specified in [7]):

$$\begin{aligned} & \sum_{i=1}^n [\text{SampleLevelRandomly}(M) \geq a] \\ &= \sum_{i=1}^n \left[\frac{-\ln(\text{unif}(0, 1))}{\ln M} \geq a \right] \\ &= \sum_{i=1}^n [\text{unif}(0, 1) \leq e^{-a \ln M}] \\ &= \sum_{i=1}^n e^{-a \ln M} = n e^{-a \ln M} \end{aligned}$$

Note: $[]$ is indicator function, $[x] = 1$ if and only if condition x is true, unif is the uniform distribution.

Then let $g(n)$ be the expected time complexity of all SearchLayer methods when the current HNSW have n objects. It's equal to the sum of the searching time on every layer.

$$\begin{aligned} g(n) &= \sum_{a=0}^{\infty} f(n e^{-a \ln M}) = \sum_{a=0}^{\infty} (n e^{-a \ln M})^c \\ &= n^c \sum_{a=0}^{\infty} (e^{-c \ln M})^a \\ &= n^c \cdot \frac{1}{1 - e^{-c \ln M}} \\ &= O(n^c) \end{aligned}$$

Naturally, we have the expression of $h(n)$ (symbol reused here for the proof of the second part):

$$\begin{aligned} h(n) &= \sum_{i=0}^{n-1} g(i) \\ &= \sum_{i=0}^{n-1} O(i^c) \leq \sum_{i=0}^{n-1} O(n^c) \\ &= O(n^{c+1}) \end{aligned}$$

Then the index time that only considers main cost (search-Layer process) is:

$$\begin{aligned} \sum_{i=0}^{t-1} 2^i h\left(\frac{n}{2^i}\right) &= \sum_{i=0}^{t-1} \left(\frac{n}{2^i}\right)^{c+1} 2^i \\ &= \sum_{i=0}^{t-1} \frac{n^{c+1}}{(2^i)^c} = h(n) \sum_{i=0}^{t-1} \frac{1}{(2^c)^i} \\ &= h(n) \cdot \frac{1 - (1/2^c)^t}{1 - (1/2^c)} \\ &= h(n) \cdot \frac{1 - (1/2^c)^{\log(n/B)+1}}{1 - 1/2^c} \end{aligned}$$

Factor $\frac{1 - (1/2^c)^{\log(n/B)+1}}{1 - 1/2^c}$ is a constant term and is bounded in the worst case $c = 1/2$:

$$\frac{1 - \frac{\sqrt{2}}{2}^{(\log(n/B)+1)}}{1 - \frac{\sqrt{2}}{2}} \leq \frac{1}{1 - \frac{\sqrt{2}}{2}}$$

We can also treat it using the master theorem (ignoring B and constants): $T(n) = 2T(n/2) + O(n^{c+1})$. By the master theorem, we get $T(n) = O(n^{c+1})$, which is consistent with our previous conclusion. $\square$

D. Proof of Theorem 5

Proof. Since in FRASP, the only extra cost to retrieve all possible edges from object o is to iterate the $N_b^{\text{del}}[o]$. Considering all n objects, the expected size of $N_b^{\text{del}}[o]$ for each object o is $O(M)$ since there are at most $O(nM)$ edges deleted in HNSW with size n and the average degree of $N_b^{\text{del}}[o]$ is at most $\frac{O(nM)}{n} = O(M)$. So the total expected time to process an object is still $O(M)$. (In our experiments, since the number of deleted edges is actually far less than the theoretical upper limit, this part actually takes very little time.) We further show that the number of edges traversed per object in one query process is $O(M)$. According to [7], in a single HNSW with multiple layers, the expected out-degree of an object over all layers is $M + \frac{1}{\ln M}M = O(M)$ in brief. Since at most three HNSW graphs are used in one query and the two child HNSWs are disjoint, each object can appear in at most two HNSWs. Thus, the total number of explored edges per object remains $O(M)$. $\square$

E. FindFirstSplitPosition

The pseudocode of FindFirstSplitPosition is shown in Algorithm 8. It starts from the root segment $[l, r]$ (Line 1) and iteratively checks if a segment has middle point mid within the query range $[L, R]$ or not. If mid is in $[L, R]$, we find the first-split position $fs = mid$, and the left and right child segments are $[l, mid]$ and $[mid + 1, r]$ respectively, and we break out of the loop (Line 4-6). If mid is larger or equal than R , we move to the left child segment $[l, mid]$ to continue (Line 7-8). Otherwise, we move to the right child segment $[mid + 1, r]$ (Line 9-10). The segment tree is built with leaf size $O(B)$, so we can guarantee that we can find the first-split position fs for any query range larger than B . The running time of the FindFirstSplitPosition is $O(\log(n/B))$.

IX. ADDITIONAL EXPERIMENTAL RESULTS

A. Experiments on Distance Relative Error

Here we report the complete results of QPS and relative error evaluation on all datasets under all query workloads, as shown in Figures 10, 11, 12, 13, and 14. Our method FRASP consistently outperforms the baseline methods in terms of both QPS and relative error under almost all settings, and often achieves a significant speedup compared with existing methods. For example, on the mixed workload of Sift, FRASP and iRangeGraph are the top two performers, while the other methods are significantly slower or less accurate. FRASP achieves a significantly better trade-off between QPS and relative error compared to iRangeGraph. The other baselines may yield a too large relative error, thus not shown in the range of the figures. The performance improvement of FRASP over existing methods is maintained when query ranges are varied.

Algorithm 8: FindFirstSplitPosition($[L, R]$)

Input: Query range $[L, R]$
Output: First-split position fs
 // Start from root segment to find first-split position
 1 $[l, r] \leftarrow [1, n];$
 2 **while** $[l, r]$ is not a leaf segment **do**
 3 $mid \leftarrow \lfloor (l + r)/2 \rfloor;$
 4 **if** $L \leq mid$ and $R > mid$ **then**
 5 $fs \leftarrow mid;$
 6 **break;**
 7 **else if** $R \leq mid$ **then**
 8 $[l, r] \leftarrow [l, mid];$
 9 **else**
 10 $[l, r] \leftarrow [mid + 1, r];$
 11 **return** $fs;$

B. Experimental Results on Total Distance Error

Figures 15, 16, 17, 18 and 19 show the total distance ratio and QPS trade off of all methods on all datasets under all query workloads. FRASP consistently outperforms the baseline methods in terms of both QPS and total distance ratio under almost all settings, and often achieves a significant speedup compared with existing methods. In these figures, we can see that FRASP achieves a significantly better trade-off between QPS and total distance ratio compared to iRangeGraph, as well as the other baselines that yield large total distance ratios. The performance improvement of FRASP over existing methods is maintained when query ranges are varied.

C. Experimental Results on Number of Distance Computations

As reported in Table IV of the experiments in Section VI-B, the total time of distance computation takes a large proportion of the total query time, but also may take a small proportion of the total time in some datasets, such as Sift and Redcaps. Here we report the number of distance computations of all methods on all datasets under the mixed workloads, as shown in Figure 20. Observe that on WIT, Gist, and rgb, FRASP achieves a significant reduction in the number of distance computations compared to the other methods, while on Sift and Redcaps, FRASP achieves a similar number of distance computations as iRangeGraph. As explained, the other computations, such as edge processing, graph exploration, and queue management, are also time-consuming on Sift and Redcaps. Our design only requires $O(M)$ expected processing time for each object, while iRangeGraph requires an amortized $O(M + \log n)$ time for each object. This explains that we can still achieve higher QPS than iRangeGraph on Sift and Redcaps, even with a similar number of distance computations.

— irange — our — superfilter — segmentgraph — ACORN — milvus




Fig. 10. QPS-Relative error on WIT.




Fig. 11. QPS-Relative error on Sift.




Fig. 12. QPS-Relative error on Gist.




Fig. 13. QPS-Relative error on YT-rgb.




Fig. 14. QPS-Relative error on Redcaps.

— FRASP — iRangeGraph — SuperPostfiltering — SeRF — ACORN — Milvus




Fig. 15. QPS-Total distance ratio on WIT.




Fig. 16. QPS-Total distance ratio on Sift.




Fig. 17. QPS-Total distance ratio on Redcaps.




Fig. 18. QPS-Total distance ratio on Gist




Fig. 19. QPS-Total distance ratio on YT-rgb




Fig. 20. Number of distance computation-Recall on mixed workload.




Fig. 21. Number of edges traversed-Recall on mixed workload.

REFERENCES

- [1] H. Li, T. N. Chan, M. L. Yiu, and N. Mamoulis, “FEXIPRO: fast and exact inner product retrieval in recommender systems,” in *SIGMOD Conference*. ACM, 2017, pp. 835–850.
- [2] M. Grohe, “word2vec, node2vec, graph2vec, x2vec: Towards a theory of vector embeddings of structured data,” in *PODS*. ACM, 2020, pp. 1–16.
- [3] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. Yih, T. Rocktäschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *NeurIPS*, 2020.
- [4] W. Li, C. Feng, D. Lian, Y. Xie, H. Liu, Y. Ge, and E. Chen, “Learning balanced tree indexes for large-scale vector retrieval,” in *KDD*. ACM, 2023, pp. 1353–1362.
- [5] H. Deguchi, T. Watanabe, Y. Matsui, M. Utiyama, H. Tanaka, and E. Sumita, “Subset retrieval nearest neighbor machine translation,” in *ACL (1)*. Association for Computational Linguistics, 2023, pp. 174–189.
- [6] N. Rekabsaz, O. Lesota, M. Schedl, J. Brassey, and C. Eickhoff, “Tripclick: The log files of a large health web search engine,” in *SIGIR*. ACM, 2021, pp. 2507–2513.
- [7] Y. A. Malkov and D. A. Yashunin, “Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs,” *IEEE transactions on pattern analysis and machine intelligence*, vol. 42, no. 4, pp. 824–836, 2018.
- [8] M. Douze, A. Guzhva, C. Deng, J. Johnson, G. Szilvasy, P.-E. Mazaré, M. Lomeli, L. Hosseini, and H. Jégou, “The faiss library,” 2024.
- [9] Y. Xu, J. Gao, Y. Gou, C. Long, and C. S. Jensen, “iRangeGraph: improvising range-dedicated graphs for range-filtering nearest neighbor search,” *Proc. ACM Manag. Data*, vol. 2, no. 6, pp. 239:1–239:26, 2024.
- [10] C. Zuo, M. Qiao, W. Zhou, F. Li, and D. Deng, “SeRF: segment graph for range-filtering approximate nearest neighbor search,” *Proc. ACM Manag. Data*, vol. 2, no. 1, pp. 69:1–69:26, 2024.
- [11] J. Wang, X. Yi, R. Guo, H. Jin, P. Xu, S. Li, X. Wang, X. Guo, C. Li, X. Xu, K. Yu, Y. Yuan, Y. Zou, J. Long, Y. Cai, Z. Li, Z. Zhang, Y. Mo, J. Gu, R. Jiang, Y. Wei, and C. Xie, “Milvus: A purpose-built vector data management system,” in *SIGMOD Conference*, 2021, pp. 2614–2627.
- [12] Q. Zhang, S. Xu, Q. Chen, G. Sui, J. Xie, Z. Cai, Y. Chen, Y. He, Y. Yang, F. Yang, M. Yang, and L. Zhou, “VBASE: unifying online vector similarity search and relational queries via relaxed monotonicity,” in *OSDI*. USENIX Association, 2023, pp. 377–395.
- [13] L. Patel, P. Kraft, C. Guestrin, and M. Zaharia, “ACORN: Performant and predicate-agnostic search over vector embeddings and structured data,” *Proceedings of the ACM on Management of Data*, vol. 2, no. 3, pp. 1–27, 2024.
- [14] Z. Wang, J. Zhang, and W. Hu, “WoW: A window-to-window incremental index for range-filtering approximate nearest neighbor search,” in *Proc. ACM SIGMOD Int. Conf. Management of Data*. New York, NY, USA: ACM, 2026, also available as arXiv:2508.18617.
- [15] A. Liang, P. Zhang, B. Yao, Z. Chen, Y. Song, and G. Cheng, “UNIFY: Unified index for range filtered approximate nearest neighbors search,” *Proc. VLDB Endow.*, vol. 18, no. 4, pp. 1118–1130, Dec. 2024. [Online]. Available: <https://doi.org/10.14778/3717755.3717770>
- [16] M. de Berg, O. Cheong, M. J. van Kreveld, and M. H. Overmars, *Computational geometry: algorithms and applications*, 3rd Edition. Springer, 2008.
- [17] Y. Xiang, J. Wang, C. Feng, and J. Shi, “FRASP: Code and Technical Report,” <https://github.com/wjy-yy/FRASP>, 2026.
- [18] Y. Malkov, A. Ponomarenko, A. Logvinov, and V. Krylov, “Approximate nearest neighbor algorithm based on navigable small world graphs,” *Information Systems*, vol. 45, pp. 61–68, 2014.
- [19] C. Fu, C. Xiang, C. Wang, and D. Cai, “Fast approximate nearest neighbor search with the navigating spreading-out graph,” *Proc. VLDB Endow.*, vol. 12, no. 5, pp. 461–474, 2019.
- [20] S. Jayaram Subramanya, F. Devvrit, H. V. Simhadri, R. Krishnawamy, and R. Kadekodi, “Diskann: Fast accurate billion-point nearest neighbor search on a single node,” *Advances in neural information processing Systems*, vol. 32, 2019.
- [21] M. Wang, X. Xu, Q. Yue, and Y. Wang, “A comprehensive survey and experimental comparison of graph-based approximate nearest neighbor search,” *Proceedings of the VLDB Endowment*, vol. 14, no. 11, pp. 1964–1978, 2021.
- [22] A. Beygelzimer, S. Kakade, and J. Langford, “Cover trees for nearest neighbor,” in *ICML*, 2006, pp. 97–104.
- [23] Y. Gu, Z. Napier, Y. Sun, and L. Wang, “Parallel cover trees and their applications,” in *Proceedings of the 34th ACM Symposium on Parallelism in Algorithms and Architectures*, 2022, pp. 259–272.
- [24] Q. Huang, J. Feng, Y. Zhang, Q. Fang, and W. Ng, “Query-aware locality-sensitive hashing for approximate nearest neighbor search,” *Proceedings of the VLDB Endowment*, vol. 9, no. 1, pp. 1–12, 2015.
- [25] P. Indyk and R. Motwani, “Approximate nearest neighbors: towards removing the curse of dimensionality,” in *Proceedings of the thirtieth annual ACM symposium on Theory of computing*, 1998, pp. 604–613.
- [26] Y. Tao, K. Yi, C. Sheng, and P. Kalnis, “Efficient and accurate nearest neighbor and closest pair search in high-dimensional space,” *ACM Transactions on Database Systems (TODS)*, vol. 35, no. 3, pp. 1–46, 2010.
- [27] Q. Chen, B. Zhao, H. Wang, M. Li, C. Liu, Z. Li, M. Yang, and J. Wang, “Spann: Highly-efficient billion-scale approximate nearest neighborhood search,” *Advances in Neural Information Processing Systems*, vol. 34, pp. 5199–5212, 2021.
- [28] C. Fu, C. Wang, and D. Cai, “High dimensional similarity search with satellite system graph: Efficiency, scalability, and unindexed query compatibility,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 44, no. 8, pp. 4139–4150, 2021.
- [29] B. Harwood and T. Drummond, “Fann: Fast approximate nearest neighbour graphs,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 5713–5722.
- [30] F. André, A. Kermarrec, and N. L. Scouarnec, “Cache locality is not enough: High-performance nearest neighbor search with product quantization fast scan,” vol. 9, no. 4, 2015, pp. 288–299.
- [31] J. Gao and C. Long, “Rabitq: quantizing high-dimensional vectors with a theoretical error bound for approximate nearest neighbor search,” *Proceedings of the ACM on Management of Data*, vol. 2, no. 3, pp. 1–27, 2024.
- [32] R. Guo, P. Sun, E. Lindgren, Q. Geng, D. Simcha, F. Chern, and S. Kumar, “Accelerating large-scale inference with anisotropic vector quantization,” in *ICML*, vol. 119, 2020, pp. 3887–3896.
- [33] H. Jégou, M. Douze, and C. Schmid, “Product quantization for nearest neighbor search,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 33, no. 1, pp. 117–128, 2011.
- [34] K. Echihabi, T. Palpanas, and K. Zoumpatianos, “New trends in high-d vector similarity search: Ai-driven, progressive, and distributed,” *Proc. VLDB Endow.*, vol. 14, no. 12, pp. 3198–3201, 2021.
- [35] W. Li, Y. Zhang, Y. Sun, W. Wang, M. Li, W. Zhang, and X. Lin, “Approximate nearest neighbor search on high dimensional data—experiments, analyses, and improvement,” *IEEE Transactions on Knowledge and Data Engineering*, vol. 32, no. 8, pp. 1475–1488, 2019.
- [36] M. Wang, L. Lv, X. Xu, Y. Wang, Q. Yue, and J. Ni, “An efficient and robust framework for approximate nearest neighbor search with attribute constraint,” in *NeurIPS*, 2023.
- [37] S. Gollapudi, N. Karia, V. Sivashankar, R. Krishnaswamy, N. Begwani, S. Raz, Y. Lin, Y. Zhang, N. Mahapatro, P. Srinivasan, A. Singh, and H. V. Simhadri, “Filtered-diskann: Graph algorithms for approximate nearest neighbor search with filters,” in *WWW*, 2023, pp. 3406–3416.
- [38] J. Mohoney, A. Pacaci, S. R. Chowdhury, A. Mousavi, I. F. Ilyas, U. F. Minhas, J. Pound, and T. Rekatsinas, “High-throughput vector similarity search in knowledge graphs,” *Proceedings of the ACM on Management of Data*, vol. 1, no. 2, pp. 1–25, 2023.
- [39] Y. Matsui, R. Hinami, and S. Satoh, “Reconfigurable inverted index,” in *ACM MM*, 2018, pp. 1715–1723.
- [40] C. Zuo and D. Deng, “Arkgraph: All-range approximate k-nearest-neighbor graph,” *Proceedings of the VLDB Endowment*, vol. 16, no. 10, pp. 2645–2658, 2023.
- [41] J. Engels, B. Landrum, S. Yu, L. Dhulipala, and J. Shun, “Approximate nearest neighbor search with window filters,” in *International Conference on Machine Learning*. PMLR, 2024, pp. 12 469–12 490.
- [42] P. Indyk and R. Motwani, “Approximate nearest neighbors: Towards removing the curse of dimensionality,” in *STOC*. ACM, 1998, pp. 604–613.
- [43] P. Indyk and H. Xu, “Worst-case performance of popular approximate nearest neighbor search implementations: Guarantees and limitations,” in *NeurIPS*, 2023.
- [44] A. Rubinstein, “Hardness of approximate nearest neighbor search,” in *STOC*. ACM, 2018, pp. 1260–1268.

- [45] M. Jiang, Z. Yang, F. Zhang, G. Hou, J. Shi, W. Zhou, F. Li, and S. Wang, "DIGRA: A dynamic graph indexing for approximate nearest neighbor search with range filter," *Proc. ACM Manag. Data*, vol. 3, no. 3, pp. 148:1–148:26, 2025.
- [46] P. E. O'Neil, E. Cheng, D. Gawlick, and E. J. O'Neil, "The log-structured merge-tree (lsm-tree)," *Acta Informatica*, vol. 33, no. 4, pp. 351–385, 1996.
- [47] M. Aumüller, E. Bernhardsson, and A. Faithfull, "Ann-benchmarks: A benchmarking tool for approximate nearest neighbor algorithms," *Information Systems*, vol. 87, p. 101374, 2020.
- [48] WIT, "WIT: Wikipedia-based Image Text Dataset," <https://github.com/google-research-datasets/wit>, 2025.
- [49] Youtube-8M, "Youtube-8M: A Large and Diverse Labeled Video Dataset for Video Understanding Research," <https://research.google.com/youtube8m/>, 2025.
- [50] S. Abu-El-Haija, N. Kothari, J. Lee, P. Natsev, G. Toderici, B. Varadarajan, and S. Vijayanarasimhan, "Youtube-8m: A large-scale video classification benchmark," *arXiv preprint arXiv:1609.08675*, 2016.
- [51] ANN-benchmarks, "ANN-Benchmarks," <http://ann-benchmarks.com/>, 2025.
- [52] Deep10M, <https://research.yandex.com/blog/benchmarks-for-billion-scale-similarity-search>, 2016.
- [53] M. Patella and P. Ciaccia, "The many facets of approximate similarity search," in *First International Workshop on Similarity Search and Applications (sisap 2008)*. IEEE, 2008, pp. 10–21.